

SQL Window Function

Deep-Dive Advanced SQL

Advanced SQL

04 — Window Functions

Torsten Grust
Universität Tübingen, Germany

1 | Window Functions

With SQL:2003, the ISO SQL Standard introduced **window functions**, a new mode of row-based computation:

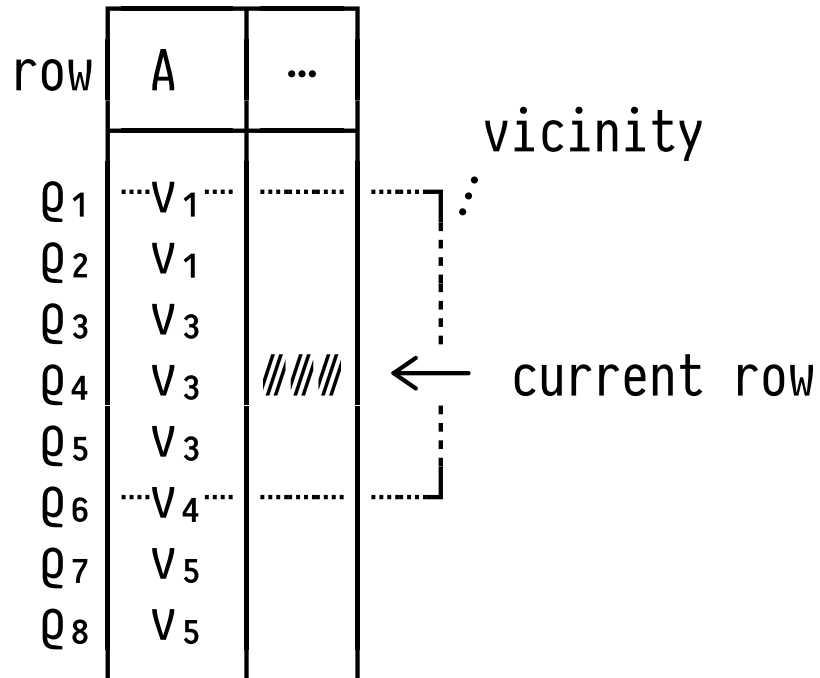
SQL Modes of Computation

SQL Feature	Mode of Computation
function	row → row
table-generating function	row → table of rows
aggregate	group of rows → row (one per group)
window function 🍷	row vicinity → row (one per row)

Window functions ...

- ... are **row-based**: each individual input row **r** is mapped to one result row,
- ... use the **vicinity** around **r** to compute this result row.

Row Vicinity: Window Frames



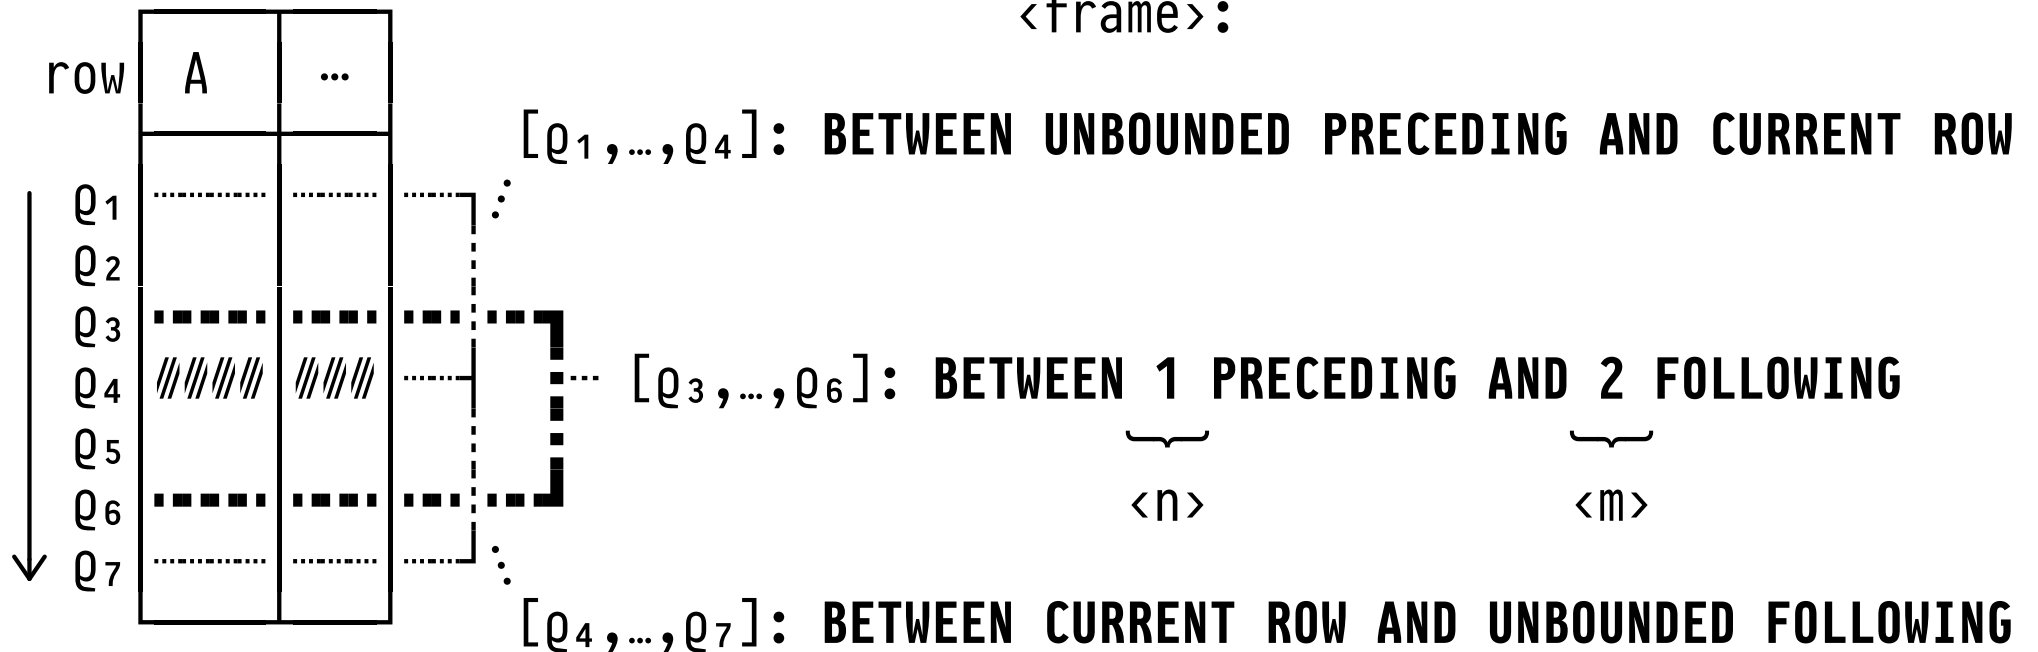
- Each row is the **current row** at one point in time.
- Row vicinity (**window, frame**) is based on either:
 - ① row **position** (**ROWS** windows)
 - ② row **values** v_i (**RANGE** windows)
- As the current row changes, the window *slides* with it.

-  Window semantics depend on a defined **row ordering**.

Window Frame Specifications (Variant: **ROWS**)

window function	ordering criteria	frame specification
$\langle f \rangle$	$\text{OVER } (\text{ORDER BY } \langle e_1 \rangle, \dots, \langle e_n \rangle$	$[\text{ ROWS } \langle \text{frame} \rangle])$

$\langle \text{frame} \rangle$:



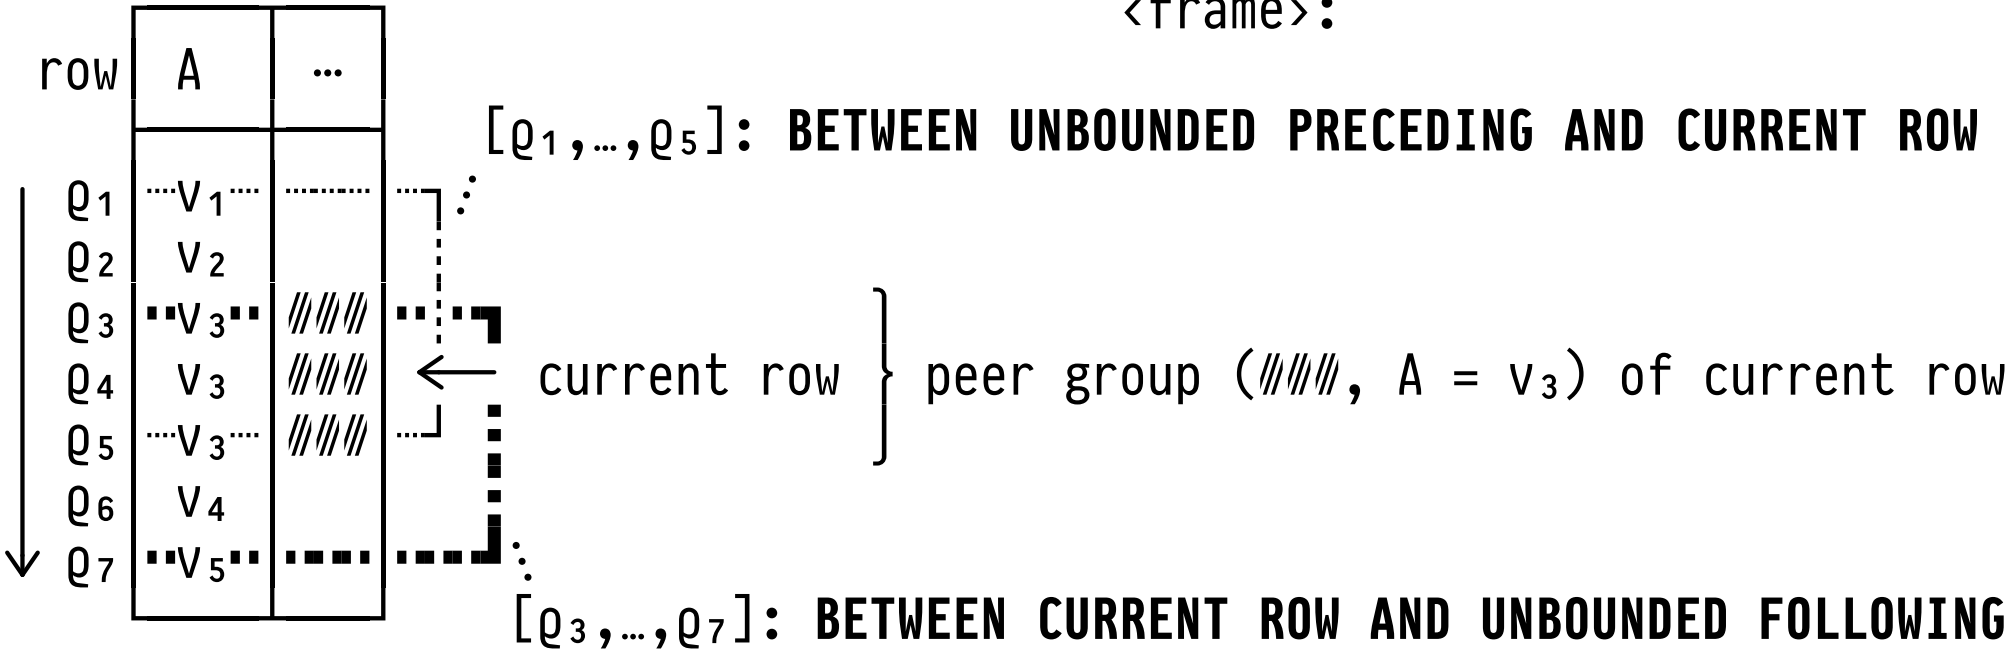
Window Frame Specifications (Variant: **RANGE**)

window function

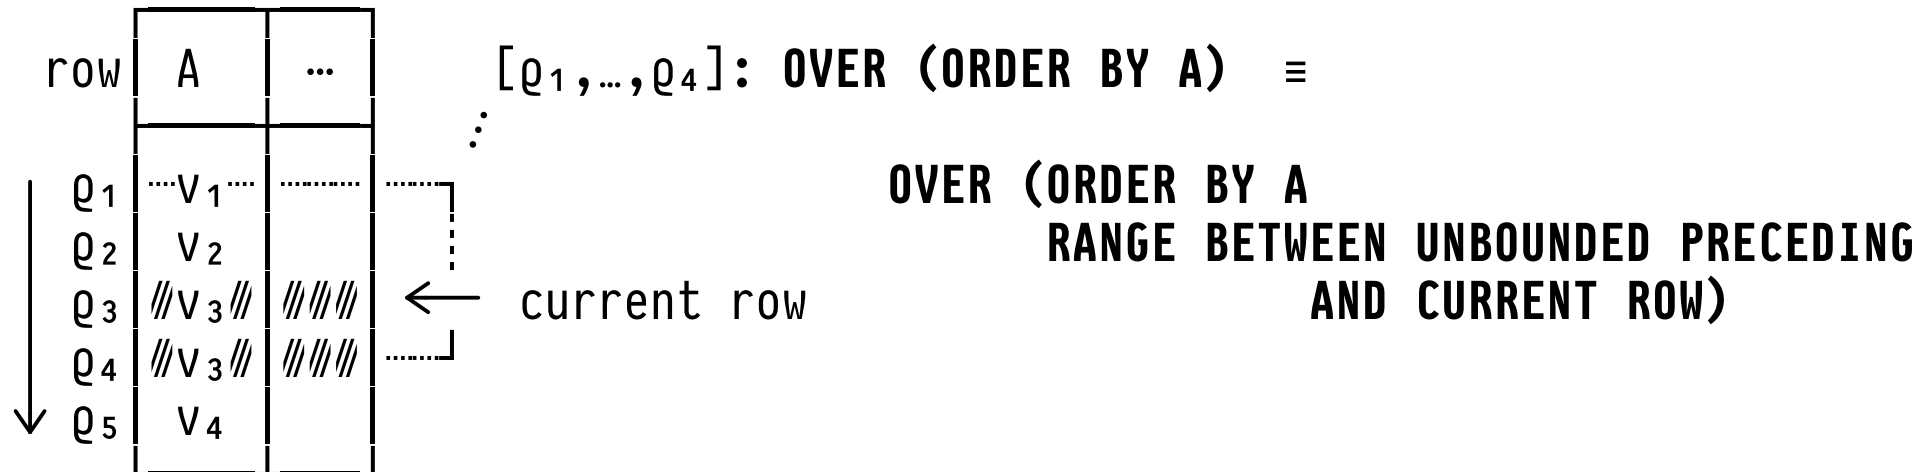
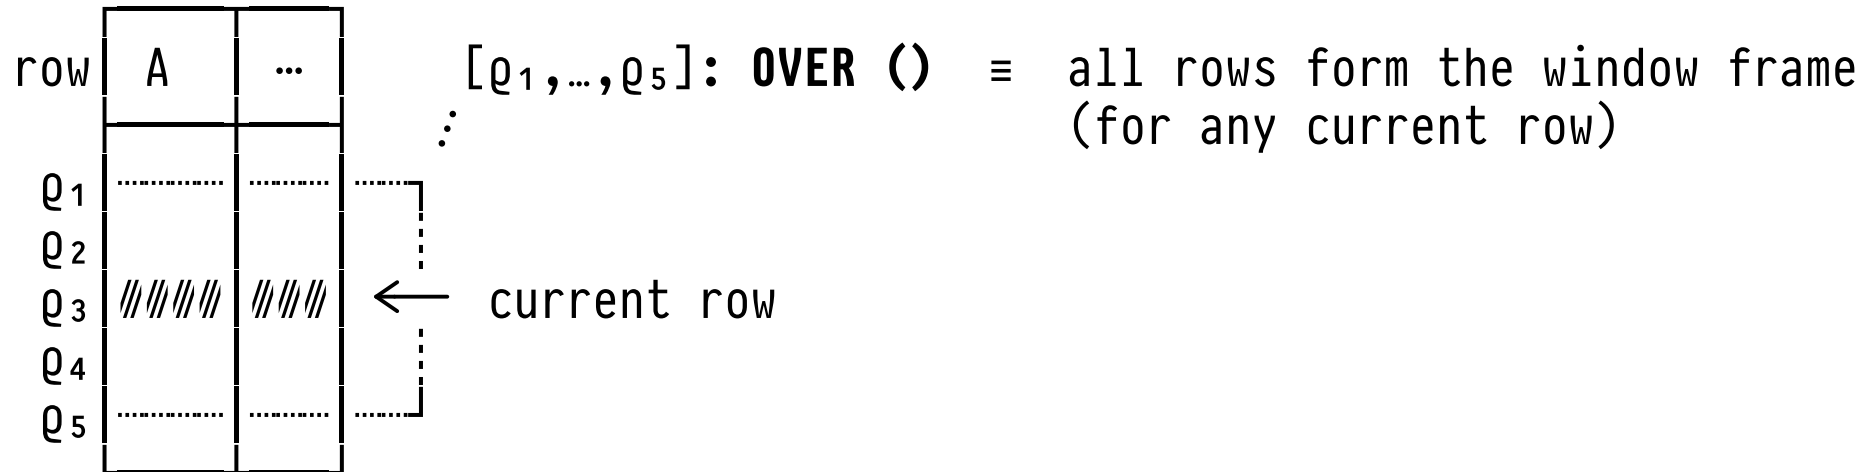
frame specification

$\underbrace{\hspace{1cm}}$
 $\langle f \rangle$ **OVER** (**ORDER BY** A [**RANGE** $\underbrace{\hspace{1cm}}$ $\langle \text{frame} \rangle$])

$\langle \text{frame} \rangle$:



Window Frame Specifications: Abbreviations



WINDOW Clause: Name the Frame

Syntactic ©: If window frame specifications

1. become unwieldy because of verbose SQL syntax and/or
2. one frame is used multiple times in a query,

add a **WINDOW** clause to a SFW block to **name the frame**, e.g.:

```
SELECT ... <f> OVER <Wi> ... <g> OVER <Wj> ...  
FROM ...  
WHERE ...  
⋮  
WINDOW <W1> AS (<frame1>), ..., <Wn> AS (<framen>)  
ORDER BY ...
```


Use SQL Itself to Explain Window Frame Semantics

Regular **aggregates** may act as window functions **<f>**. All **rows in the frame will be aggregated**:

```
SELECT w.row AS "current row",  
       COUNT(*) OVER win AS "frame size",  
       array_agg(w.row) OVER win AS "rows in frame"  
FROM   W AS w  
WINDOW win AS (<frame>)
```

Table **W**

<u>row</u>	a	b
q ₁	1	●
q ₂	2	○
q ₃	3	○
q ₄	3	●
⋮	⋮	⋮

2 | **PARTITION BY:** Window Frames Inside Partitions

Optionally, we may **partition** the input table *before* rows are sorted and window frames are determined:

all input rows that agree on all $\langle p_i \rangle$ form one partition

```
<f> OVER ( [ PARTITION BY  $\langle p_1 \rangle, \dots, \langle p_m \rangle$  ]  
           [ ORDER BY  $\langle e_1 \rangle, \dots, \langle e_n \rangle$  ]  
           [ <frame> ] )
```

- Note:

1. Frames **never cross partitions.**
2. **BETWEEN ... PRECEDING AND ... FOLLOWING** respects partition boundaries.

Y Q: What is the Chance of Fine Weather on Weekends?

Input: Daily weather readings in **sensors**:


Table **sensors**

<u>day</u>	weekday	temp	rain
1	Fri	10	800
2	Sat	12	300
⋮	⋮	⋮	⋮

- The weather is fine on day *d* if—on *d* and the two days **prior**—the minimum temperature is above 15°C and the overall rainfall is less than 600ml/m².
- **Expected output:**

weekend?	% fine
f	29
t	43

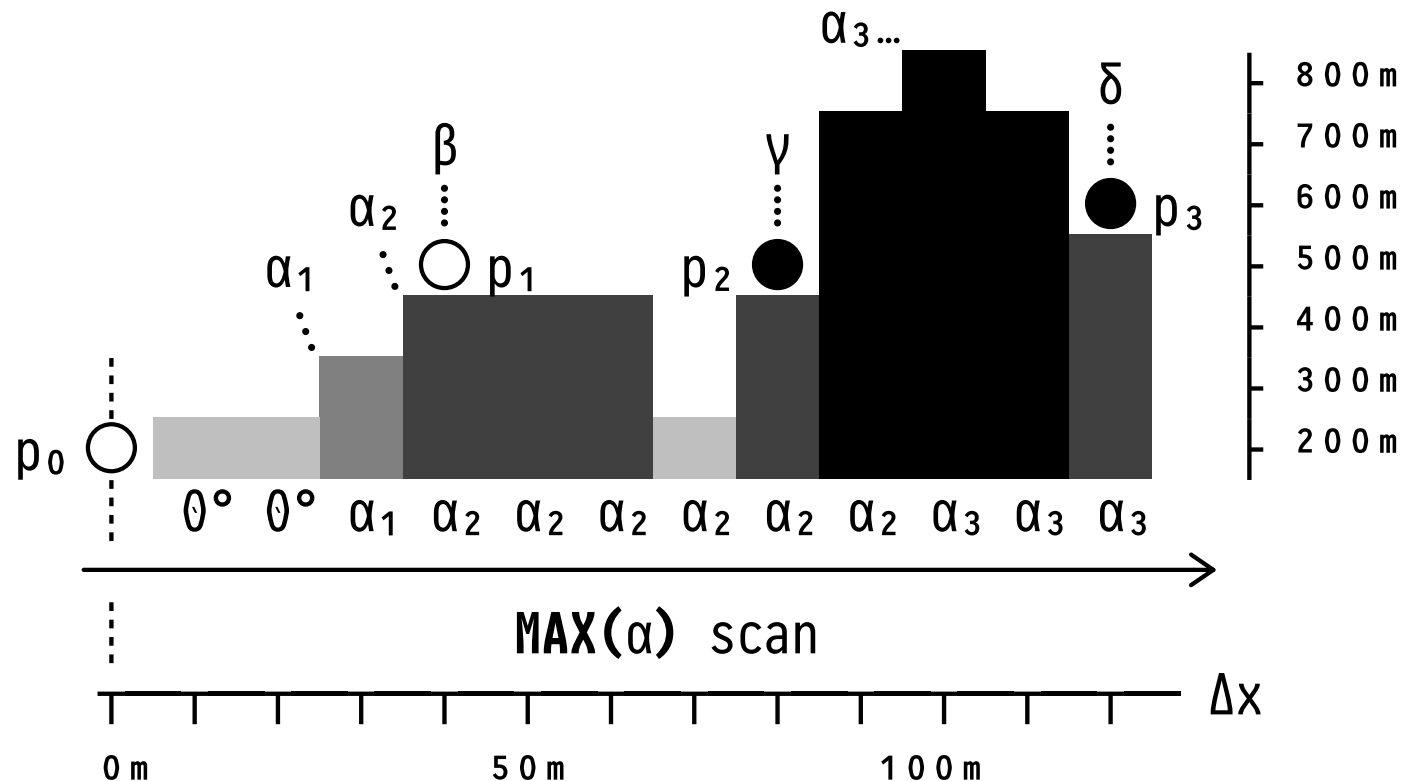

Y Q: What is Visible in a Hilly Landscape?



- From the viewpoint of p_0 (☺) we can see p_1 , but...
 - ... p_2 is **obscured** (no straight-line view from p_0),
 - ... p_3 is **obscured** (lies behind the 800m peak).

Y Q: What is Visible in a Hilly Landscape? — A: MAX Scan!

$$\alpha = \text{atan}(\Delta\text{alt}/\Delta x)$$



- We have $0^\circ < \alpha_1 < \alpha_2 < \alpha_3$ and $\beta \geq \alpha_2$, $\gamma < \alpha_2$, $\delta < \alpha_3$.
- \uparrow \uparrow \uparrow
 p_1 visible $p_{2,3}$ obscured

Y Q: What is Visible in a Hilly Landscape?

- **Input:** Location of p_0 (here: $x = 0$) and 1D-map of hills:

Table `map`

<code>x</code>	<code>alt</code>
0	200
10	200
\vdots	\vdots
120	500

- **Output:** Can p_0 see the point on the hilltop at x ?

<code>x</code>	<code>visible?</code>
0	true
10	true
\vdots	\vdots
120	false

Q: What is Visible in a Hilly Landscape? — MAX Scan

WITH

-- ❶ Angles α (in $^\circ$) between p_0 and the hilltop at x

angles(x , angle) **AS** (

SELECT $m.x$,

degrees(**atan**(($m.alt - p_0.alt$) /

abs($p_0.x - m.x$))) **AS** angle

FROM map **AS** m

WHERE $m.x > p_0.x$),

-- ❷ **MAX**(α) scan (to the right of p_0)

max_scan(x , max_angle) **AS** (

SELECT $a.x$,

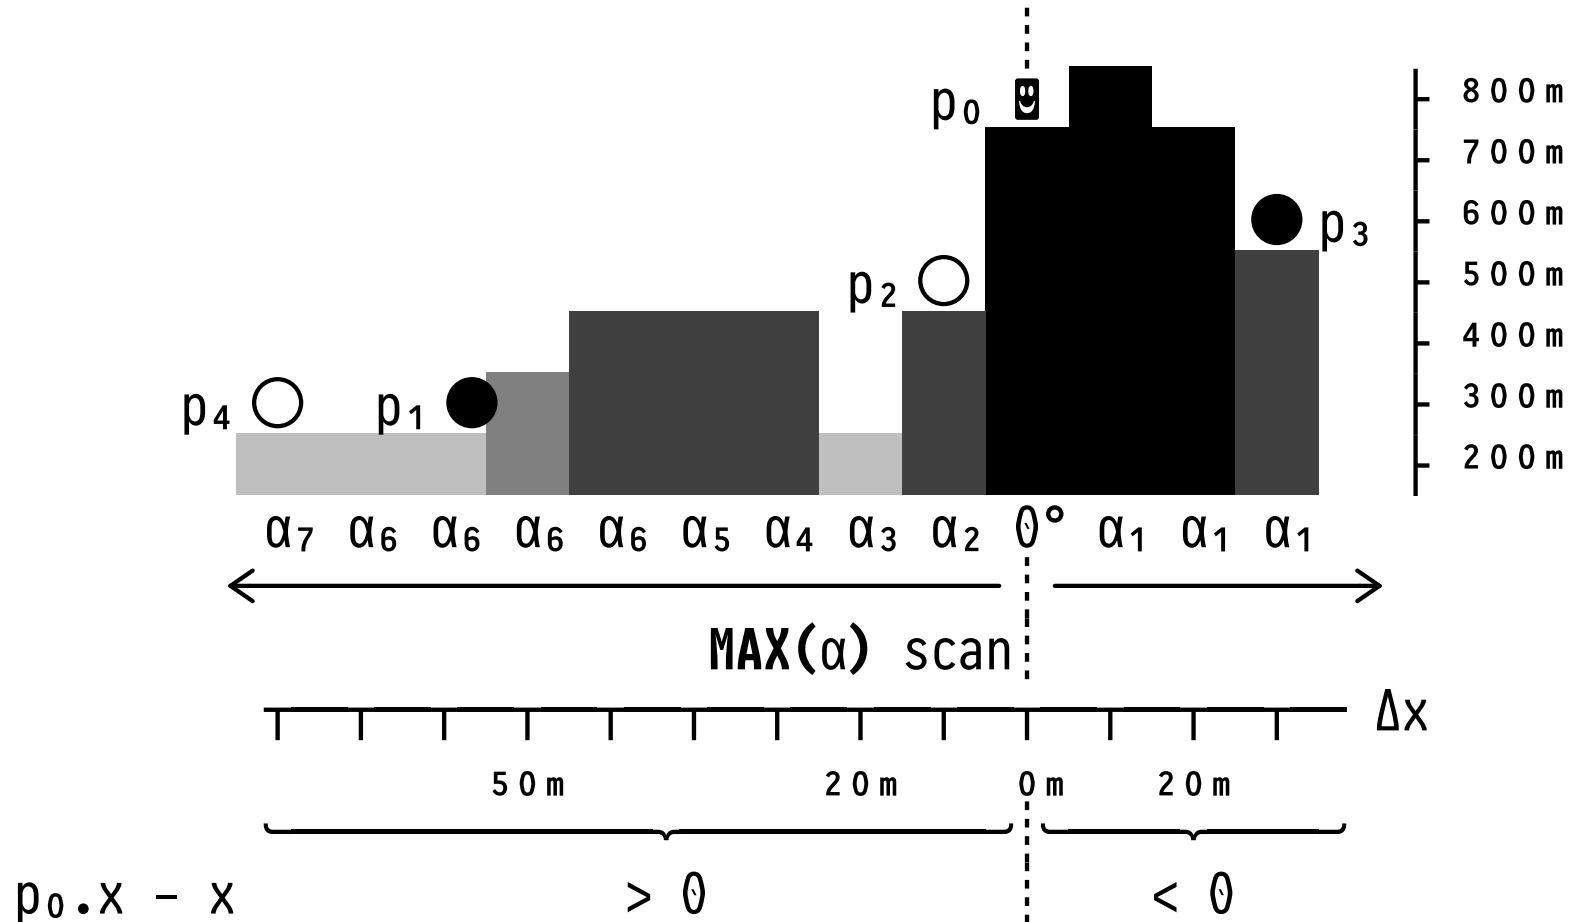
MAX($a.angle$)

OVER (**ORDER BY** $abs(p_0.x - a.x)$) **AS** max_angle

FROM angles **AS** a),

⋮

Looking Left *and* Right: PARTITION BY



- Need MAX scans left *and* right of $p_0 \Rightarrow$ use PARTITION BY.

Y Looking Left *and* Right: **PARTITION BY**

```
WITH
:
-- 2 MAX( $\alpha$ ) scan (left/right of  $p_0$ )
max_scan(x, max_angle) AS (
  SELECT a.x,          --  $\in \{-1, 0, 1\}$ 
         MAX(a.angle)   --  $\underbrace{\hspace{10em}}$ 
         OVER (PARTITION BY sign( $p_0.x - a.x$ )
              ORDER BY abs( $p_0.x - a.x$ )) AS max_angle
  FROM   angles AS a    --  $\underbrace{\hspace{10em}}$ 
                      --  $\Delta x > 0$ 
),
:
```

- $\forall a \in \text{angles}: a.x \neq p_0.x \Rightarrow$ We end up with **two** partitions.

3 | Scans: Not Only in the Hills

Scans are a general and expressive computational pattern:

$\underbrace{\langle \text{agg} \rangle(\langle e \rangle)}_{(\phi, z, \oplus)} \text{ OVER } (\text{ORDER BY } \langle e_1 \rangle, \dots, \langle e_n \rangle \{ \text{RANGE}, \text{ROWS} \} \text{ BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW})$

- Available in a variety of forms in programming languages
 - Haskell: `scanl z \oplus xs`, APL: `$\oplus \backslash xs$` , Python: `accumulate`:
`scanl \oplus z [x1, x2, ...] = [z, z \oplus x1, (z \oplus x1) \oplus x2, ...]`
- In parallel programming: *prefix sums* (👉 Guy Blelloch)
 - Sorting, lexical analysis, tree operations, reg.exp. search, drawing operations, image processing, ...

4 | Interlude: Quiz

Q: Assume $xs \equiv '((b*2)-4*a*c)*0.5'$. What is computed below?

```
SELECT inp.pos, inp.c,  
       SUM((array[1,-1,0])[COALESCE(p.oc, 3)])  
         OVER (ORDER BY inp.pos) AS d  
FROM   unnest(string_to_array(xs, NULL))  
       WITH ORDINALITY AS inp(c,pos),  
       LATERAL (VALUES (array_position(array['(', ')'],  
                                       inp.c))) AS p(oc)  
ORDER BY inp.pos;
```

💡 **Hint** (this is the same query expressed in APL):

```
xs ← '((b*2)-4*a*c)*0.5'  
+ \ (1 -1 0) ['(') ⍷ xs
```

5 | Beyond Aggregation: Window Functions

window function



```
<f> OVER ([ PARTITION BY <p1>, ..., <pm> ]  
          [ ORDER BY <e1>, ..., <en> ]  
          [ <frame> ])
```

Kinds of window functions <f>:

1. **Aggregates:** `SUM(·)`, `AVG(·)`, `MAX(·)`, `array_agg(·)`, ... ✓
2. **Row Access:** access row by *absolute/relative position* in ordered frame or partition: first/last/ n^{th} / n rows away
3. **Row Ranking:** assign numeric *rank of row* in its partition

6 | **LAG/LEAD**: Access Rows of the Past and Future

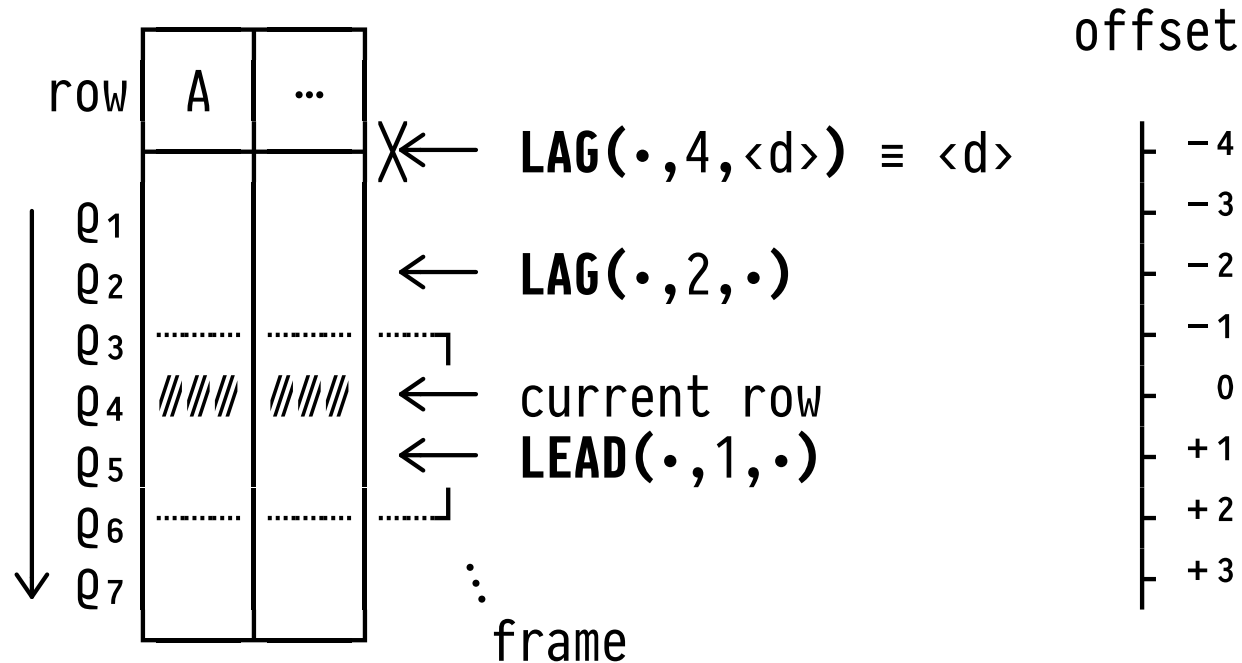
Row access at offset $\mp\langle n \rangle$, relative to the current row:

```
-- evaluate <e> as if we were  
-- <n> rows before the current row  
--  
    LAG(<e>,<n>,<d>) OVER ([ PARTITION BY <p1>,...,<pm> ]  
                           ORDER BY <e1>,...,<en>  
                           [ <frame> ])
```

Note:

- **LEAD**(<e>,<n>,<d>): ... <n> rows **after** the current row ...
- Scope is partition—may access rows outside the <frame>.
- No row at offset $\mp\langle n \rangle \Rightarrow$ return default <d>.

LAG/LEAD: Row Offsets



- The frame of the current row is irrelevant for **LAG/LEAD**.
- If no default value **<d>** given \Rightarrow return **NULL**.

Y A March Through the Hills: Ascent or Descent?

```
SELECT m.x, m.alt,  
       CASE sign(LEAD(m.alt, 1) OVER rightwards - m.alt)  
         WHEN -1 THEN '↘' WHEN 1 THEN '↗'  
         WHEN 0 THEN '→' ELSE '?'  
       END AS climb,  
       LEAD(m.alt, 1) OVER rightwards - m.alt AS "by [m]"  
FROM   map AS m  
WINDOW rightwards AS (ORDER BY m.x) -- marching right
```

x	alt	climb	by [m]
0	200	→	0
⋮	⋮	⋮	⋮
90	700	↗	100
100	800	↘	-100
110	700	↘	-200
120	500	?	NULL

Y Crime Scene: Sessionization

A spy broke into the Police HQ computer system. A `log` records keyboard activity of user `uid` at time `ts`:

Table `log`


<u>uid</u>	<u>ts</u>
00:00:00:00:00:00	09-17-2016 07:25:12
	09-17-2016 07:25:18
	09-17-2016 08:01:55
	09-17-2016 08:02:07
	09-17-2016 08:05:30
	09-17-2016 08:05:39
	09-17-2016 08:05:46

- **Q:** Can we **sessionize** the log so that investigators can identify *sessions* (\equiv streaks of uninterrupted activity)?

Y Sessionization (Query Plan)

1. Cop and spy sessions happen independently (even if concurrent): partition table **log** into 🕵️ and 🕶️ rows.
2. **Tag** keyboard activities (here: 🕵️):

$t_2 - t_1 \leq n \Rightarrow$ continue session (tag t_2 with **0**)



$t_6 - t_2 > n \Rightarrow$ new session (tag t_6 with **1**)

3. **Scan** the tagged table and derive session IDs by maintaining a **running sum** of *start of session* tags.

Y Sessionization (Query Plan)

1

uid	ts
u ₁	t ₁
u ₁	t ₂
u ₂	t ₃
u ₂	t ₄
u ₂	t ₅
u ₁	t ₆
u ₂	t ₇

2

uid	ts
u ₁	↓ t ₁
u ₁	↓ t ₂
u ₁	↓ t ₆

u ₂	↓ t ₃
u ₂	↓ t ₄
u ₂	↓ t ₅
u ₂	↓ t ₇

3

uid	ts	sos
u ₁	↓ t ₁	1
u ₁	↓ t ₂	0
u ₁	↓ t ₆	1

u ₂	↓ t ₃	1
u ₂	↓ t ₄	1
u ₂	↓ t ₅	0
u ₂	↓ t ₇	0

← log start

└─ t₆ - t₂ > n

└─ ⇒ new session

← log start

└─ t₇ - t₅ ≤ n

└─ ⇒ continue

4

uid	ts	session
u ₁	t ₁	1
u ₁	t ₂	1
u ₁	t ₆	2

u ₂	t ₃	1
u ₂	t ₄	2
u ₂	t ₅	2
u ₂	t ₇	2

- At log start, always begin a new session.
- How to assign *global session IDs* (u₂'s sessions: 3, 4)?

Y Image Compression by Run-Length Encoding

Compress image by identifying pixel **runs** of the same color:

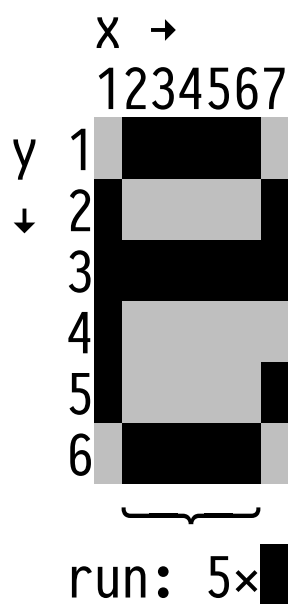


Table **original**

x	y	pixel
1	1	gray
2	1	black
⋮	⋮	⋮
6	6	black
7	6	gray










Table **encoding**


run	length	pixel
1	1	gray
2	5	black
⋮	⋮	⋮
12	5	black
13	1	gray

- Here: assumes a row-wise linearization of the pixel map.
- In b/w images we may omit column **pixel** in table **encoding**.








Y Run-Length Encoding (Query Plan)

1

x	y	pixel	change?
1	1		t 1
2	1		t 1
3	1		f 0
4	1		f 0
5	1		f 0
6	1		f 0
7	1		t 1
⋮	⋮	⋮	⋮



2

x	y	pixel	change?	Σ change?
1	1		1	1
2	1		1	2
3	1		0	2
4	1		0	2
5	1		0	2
6	1		0	2
7	1		1	3
⋮	⋮	⋮	⋮	⋮



... run #2 of length 5

- ①: `LAG(pixel,1,undefined)`: pixel @ (1,1) always “changes.”
- ②: `SUM()` scan of `change?` may serve as run identifier.

7 | **FIRST_VALUE, LAST_VALUE, NTH_VALUE**: In-Frame Row Access

Aggregates reduce *all rows* inside a frame to a single value.
Now for something different:

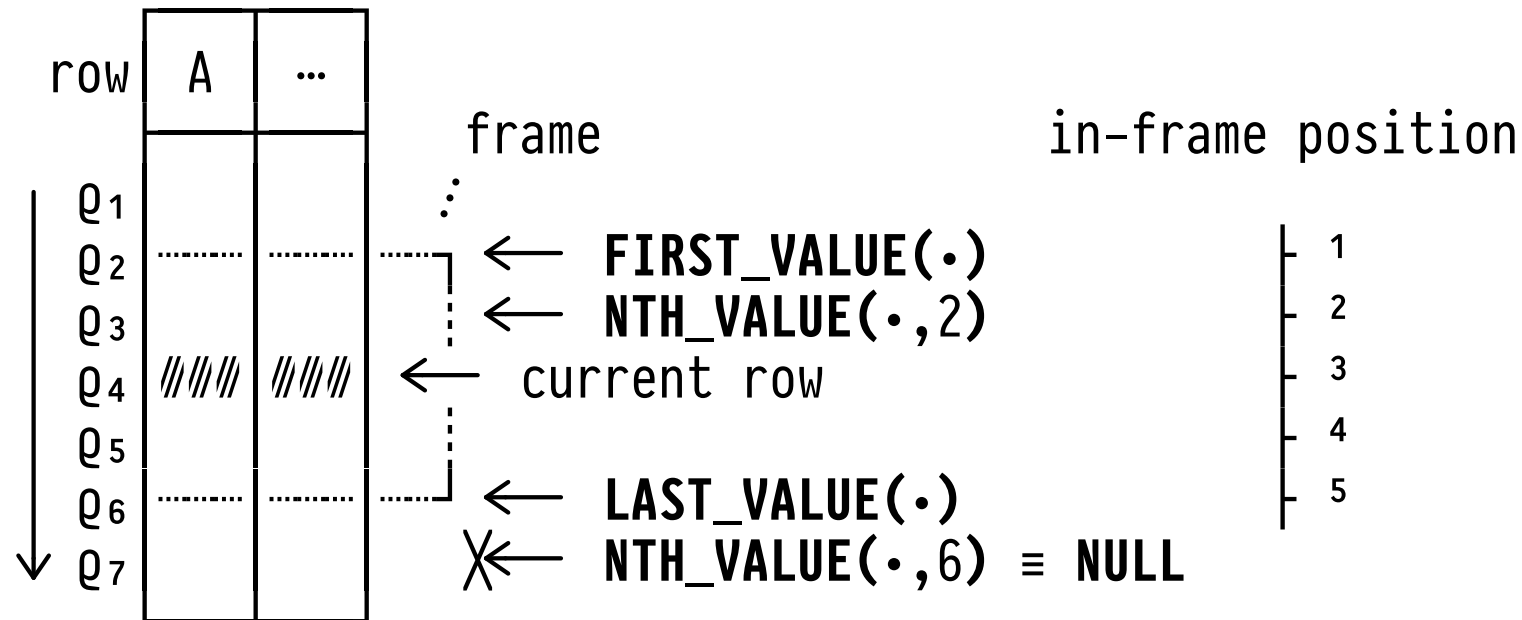
- **Positional access to individual rows** inside a frame is provided by three window functions:

```
-- evaluate expression <e> as if we were at  
-- the first/last/<n>th row in the frame  
--
```

```
    FIRST_VALUE(<e>)  
    LAST_VALUE(<e>)  
    NTH_VALUE(<e>, <n>) } OVER (...)
```

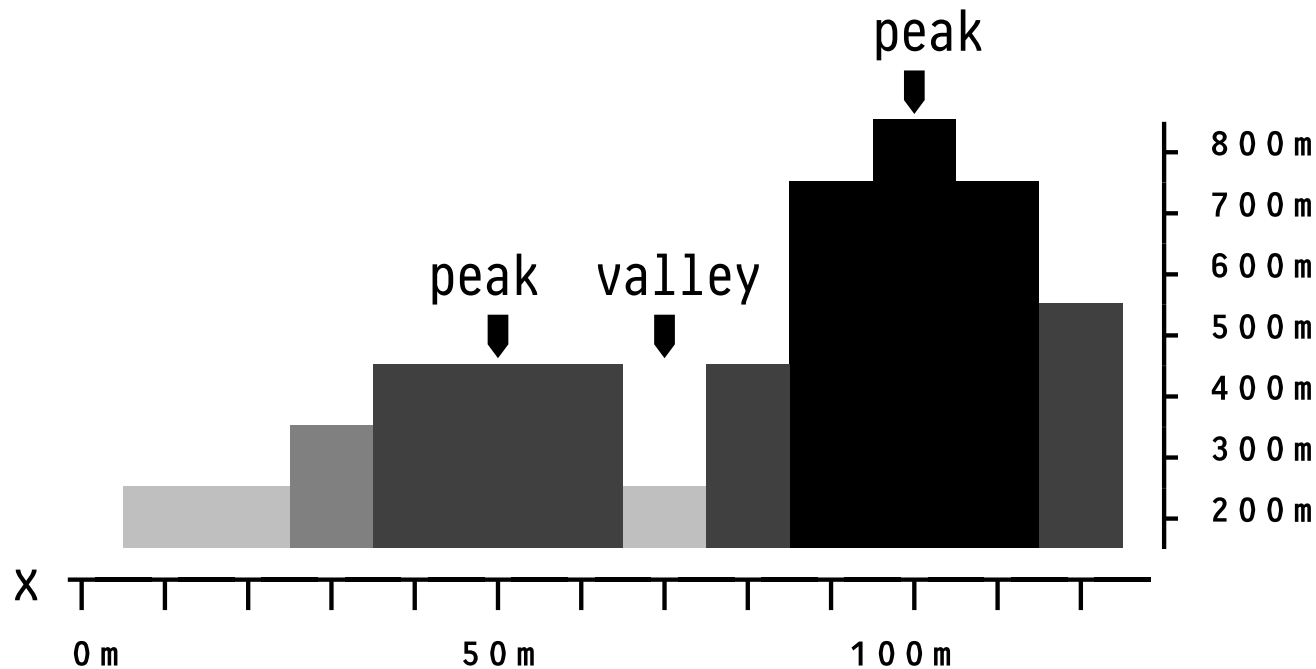
- **NTH_VALUE(<e>, <n>)**: No <n>th row in frame \Rightarrow return **NULL**.

In-Frame Row Access



- `FIRST_VALUE(<e>)` \equiv `NTH_VALUE(<e>,1)`.

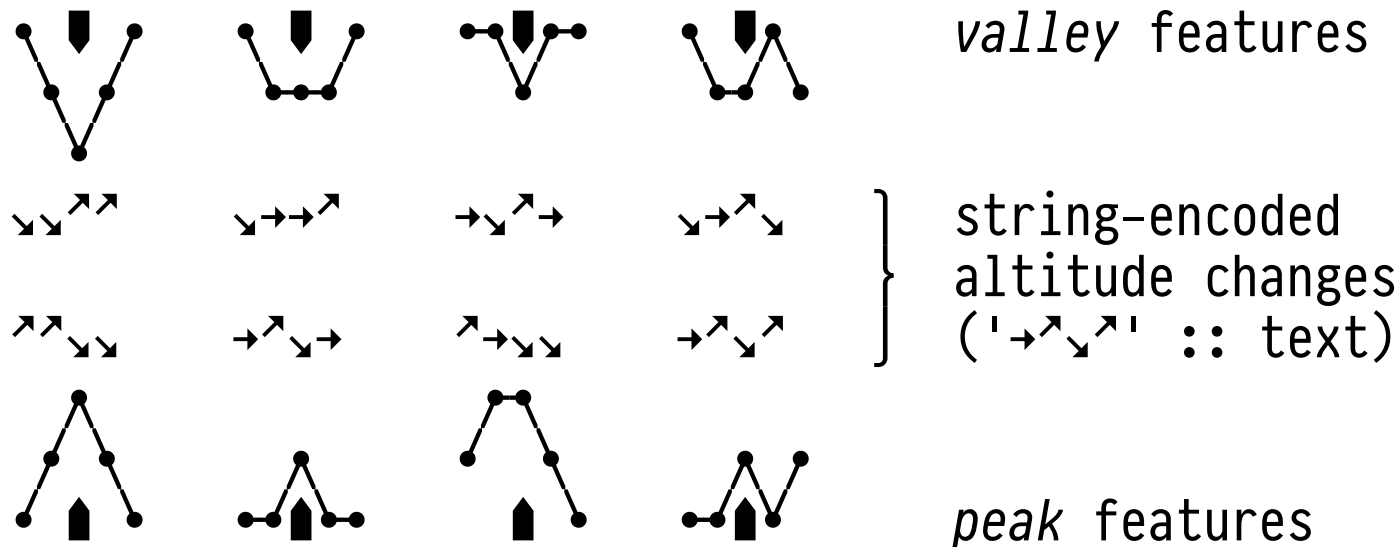
Y Detecting Landscape Features



- Detect features in hilly landscape. Attach label $\in \{\text{peak}, \text{valley}, -\}$ to every location x .
- Feature defined by relative altitude change **in vicinity**.

Y Detecting Landscape Features (Query Plan)

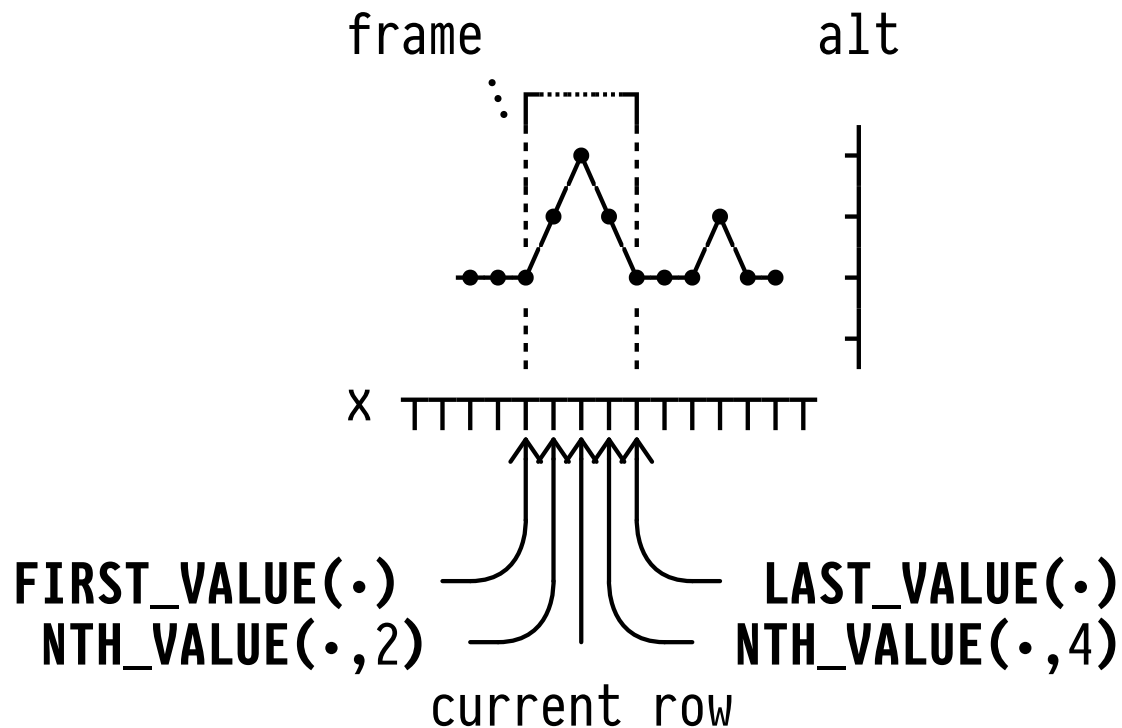
1. Track relative altitude changes in a sliding **x-window** of size 5:



2. **Pattern match** on change strings to detect features.

Y Altitude Changes in a Sliding Window

- Frame: ROWS BETWEEN 2 PRECEDING AND 2 FOLLOWING (5 rows):



- $\text{FIRST_VALUE}(\text{alt}) < \text{NTH_VALUE}(\text{alt}, 2) \Rightarrow \text{ascent ('↗')}.$



Y Altitude Changes in a Sliding Window

```
-- Find slopes in -2/+2 vicinity around point x
SELECT m.x, slope(
  sign(FIRST_VALUE(m.alt) OVER w-NTH_VALUE(m.alt,2) OVER w),
  sign(NTH_VALUE(m.alt,2) OVER w-m.alt
  sign(m.alt -NTH_VALUE(m.alt,4) OVER w),
  sign(NTH_VALUE(m.alt,4) OVER w-LAST_VALUE(m.alt) OVER w)
)
FROM   map AS m
WINDOW w AS (ORDER BY m.x ROWS BETWEEN 2 PRECEDING
                                   AND 2 FOLLOWING)
```

- Recall: 1D landscape represented as table `map(x,alt)`.
- UDF encodes altitude changes: `slope(-1,-1,0,1) ≡ '↘↘→↗'`.

Row Pattern Matching (SQL:2016)

SQL:2016 introduced an entirely new SQL construct, **row pattern matching** (`MATCH_RECOGNIZE`):

1. `ORDER BY`: Order the rows of a table.
 2. `DEFINE`: Tag rows that satisfy given predicates.
 3. `PATTERN`: Specify a **regular expression over row tags**, find matches in the ordered sequence of rows.
 4. `MEASURES`: For each match, evaluate expressions that measure its features (matched rows, length, ...).
-  As of June 2017, not supported by . Implemented in Oracle® 12i only.

Row Pattern Matching (SQL:2016)

```
SELECT *
FROM map
MATCH_RECOGNIZE (
  ORDER BY x
  MEASURES FIRST(x,1)      AS x,
               MATCH_NUMBER() AS feature,
               CLASSIFIER()  AS slope
  ONE ROW PER MATCH
  AFTER MATCH SKIP TO NEXT ROW
  PATTERN (((DOWN DOWN|DOWN EVEN|UP DOWN|EVEN DOWN)...)
  DEFINE UP      AS UP.alt > PREV(UP.alt),      --
         DOWN    AS DOWN.alt < PREV(DOWN.alt),  -- } row tags
         EVEN    AS EVEN.alt = PREV(EVEN.alt)   --
)

```

Output

x	feature	slope
50	1	DOWN
70	2	UP
100	3	DOWN

8 : Numbering and Ranking Rows

Countless problem scenarios involve the **number** (position) or **rank** of the current row in an *ordered sequence* of rows.

- Family of window functions to number/rank rows:

ROW_NUMBER()	} OVER ([PARTITION BY <p ₁ >, ..., <p _m >] [ORDER BY <e ₁ >, ..., <e _n >])	-- intra-partition ranking ✓
DENSE_RANK()		--
RANK()		--
PERCENT_RANK()		--
CUME_DIST()		-- ranking w/o ORDER BY ⚡
NTILE(<n>)		--

- Scope is partition (if present) — `<frame>` is irrelevant.

Numbering and Ranking Rows — `<f> OVER (ORDER BY A)`

Table W

row	A	ROW_NUMBER	DENSE_RANK	RANK
q1	1	1	1	1
q2	2	2	2	2
q3	3	3	3	3
q4	3	4	3	3
q5	3	5	3	3
q6	4	6	4	6
q7	5	7	5	7
q8	5	8	5	7
q9	6	9	6	9

... rows that agree on
... the sort criterion
... (here: **A**) rank equally

I mind the ranking gap
(think Olympics)

- $\text{DENSE_RANK}() \leq \text{RANK}() \leq \text{ROW_NUMBER}()$

Y Once More: Find the Top n Rows in a Group

Table `dinosaurs`

species	length	height	legs
:	:	:	$\in \{2, 4, \text{NULL}\}$

```
SELECT tallest.legs, tallest.species, tallest.height
FROM (SELECT d.legs, d.species, d.height,
ROW_NUMBER()...RANK() OVER (PARTITION BY d.legs
                                ORDER BY d.height DESC) AS n
      FROM dinosaurs AS d
      WHERE d.legs IS NOT NULL) AS tallest
WHERE n <= 3
```

- `RANK()` vs `ROW_NUMBER()`: both OK, but different semantics!
- Need a subquery: window functions *not* allowed in `WHERE`.

Y Identify Consecutive Ranges

- What you often encounter in scientific papers 🙄:
“... as Knuth has shown in [5,2,14,3,1,42,6,10,7,13] ...”
- What you want to see 😊:
“... as Knuth has shown in [1–3,5–7,10,13&14,42] ...”

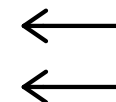
Table **citations**

ref
5
2
⋮
13



Output

ref	range
1	r_0
2	r_0
⋮	⋮
42	r_4



references belong
to the same range

Y Identify Consecutive Ranges (Query Plan)

1	2	ROW_NUMBER()					
ref	ref						
5	1	-	1	=	0	}	range 0 ≡ r ₀
2	2	-	2	=	0		
14	3	-	3	=	0		
3	5	-	4	=	1	}	range 1 ≡ r ₁
1	6	-	5	=	1		
42	7	-	6	=	1		
6	10	-	7	=	3	}	range 3 ≡ r ₂
10	13	-	8	=	5		range 5 ≡ r ₃
7	14	-	9	=	5	}	
13	42	-	10	=	32	}	range 32 ≡ r ₄
		⋮					
		subtract					

Numbering and Ranking Rows — `<f> OVER (ORDER BY A)`

row	A	PERCENT_RANK	CUME_DIST	NTILE(3)
q1	1	0	1/9	1
q2	2	1/8	2/9	1
q3	3	2/8	5/9	1
q4	3	2/8	5/9	2
q5	3	2/8	5/9	2
q6	4	5/8	6/9	2
q7	5	6/8	8/9	3
q8	5	6/8	8/9	3
q9	6	8/8	9/9	3

... rows that agree on
... the sort criterion
... (here: **A**) rank equally

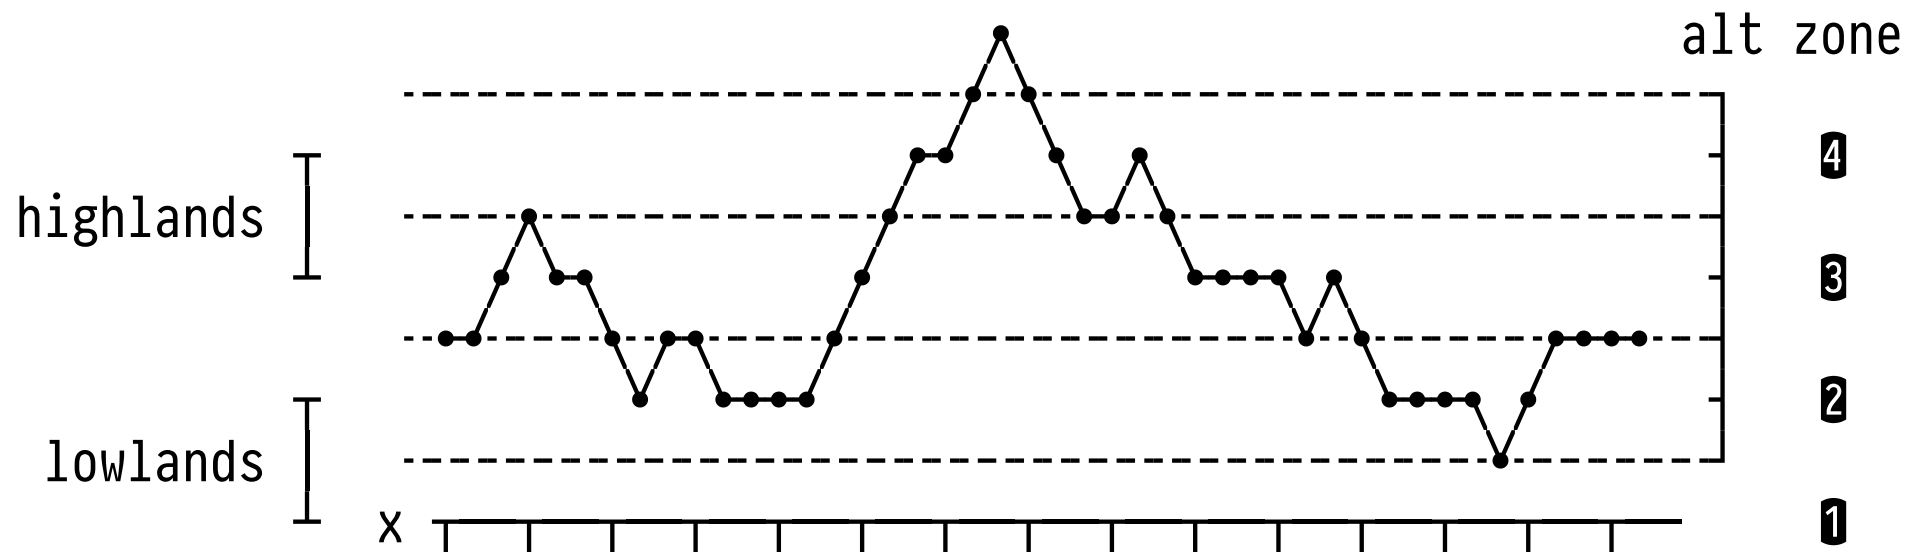
← current row is in the n^{th}
of 3 chunks of rows

$n\%$ of the other rows rank
lower than the current row

the current row and lower ranked
rows make up $n\%$ of all rows

Y Altitudinal Mountain Zones

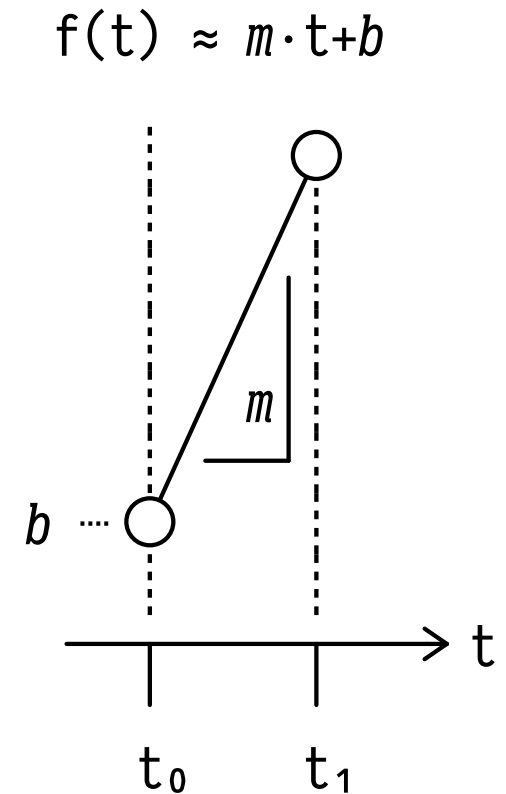
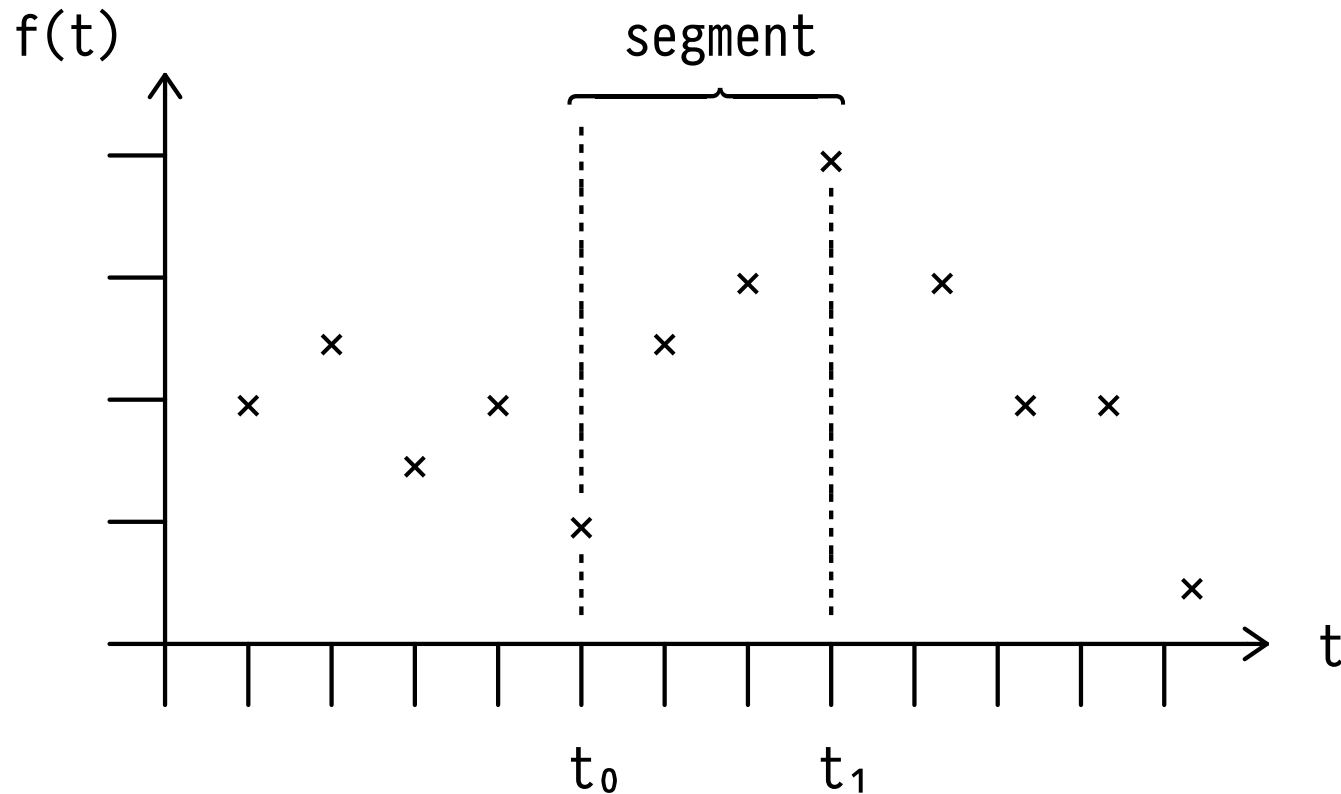
- Classify the altitudes of a mountain range into
 1. **equal-sized** vegetation **zones** and
 2. lowlands (altitude in the lowest **20%**) and highlands (between **60%–80%** of maximum altitude).



Y Altitudinal Mountain Zones

```
-- Classify altitudinal zones in table mountains(x, alt)
--
SELECT
  m.x, m.alt,
  NTILE(4) OVER altitude AS zone,
  CASE
    WHEN PERCENT_RANK() OVER altitude BETWEEN 0.6 AND 0.8
      THEN 'highlands'
    WHEN PERCENT_RANK() OVER altitude < 0.2
      THEN 'lowlands'
    ELSE '-'
  END AS region
FROM   mountains AS m
WINDOW altitude AS (ORDER BY m.alt)
ORDER BY m.x;
```

Y Linear Approximation of a Time Series



1. **NTILE(<n>)** segments time series at desired granularity.
2. Compute m , b in each **segment** \equiv **window frame**.

9 | Summary: Window Function Semantics¹

Scope	Computation	Function	Description
frame	aggregation	(aggregates)	<code>SUM</code> , <code>AVG</code> , <code>MAX</code> , <code>array_agg</code> , ...
	row access	<code>FIRST_VALUE(e)</code>	<code>e</code> at first row in frame
		<code>LAST_VALUE(e)</code>	<code>e</code> at last row in frame
		<code>NTH_VALUE(e,n)</code>	<code>e</code> at n^{th} row in frame
partition	row access	<code>LAG(e,n,d)</code>	<code>e</code> at n rows <i>before</i> current row
		<code>LEAD(e,n,d)</code>	<code>e</code> at n rows <i>after</i> current row
	ranking	<code>ROW_NUMBER()</code>	number of current row
		<code>RANK()</code>	rank with gaps (“Olympics”)
		<code>DENSE_RANK()</code>	rank without gaps
		<code>PERCENT_RANK()</code>	relative rank of current row
		<code>CUME_DIST()</code>	ratio of rows up to “—”
		<code>NTILE(n)</code>	rank on a scale $\{1,2,\dots,n\}$

¹ `FIRST_VALUE(e)`: expression `e` will be evaluated as if we are at the first row in the frame.
`LAG(e,n,d)`: default expression `d` is returned if there is no row at offset n before the current row.